CUDAHeat / GPU PIPELINE BENCHMARK

GPU pipeline scaling

Compute, device, and complete-pipeline time · full hourly record

GPU-ONLY MEASUREMENTS

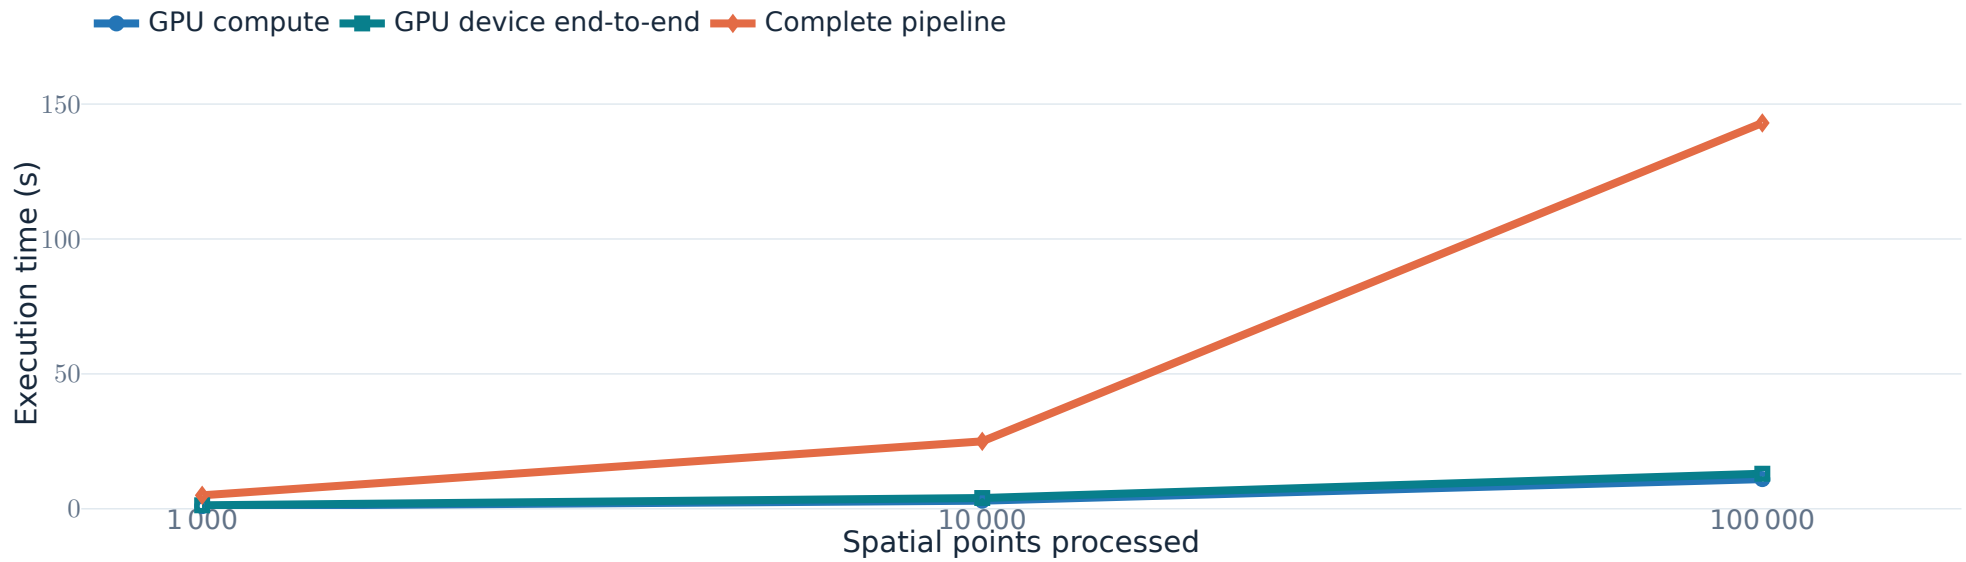

THE TAKEAWAY Host I/O and transfers remain visible beside GPU kernel time.

Source: Illustrative data only · measured GPU benchmark results

A reusable presentation theme

The same canvas, palette, font, and visual hierarchy as the Python figures

12,000	143.3 s	10.6 s	92%	99.999%
SPATIAL POINTS	FULL PIPELINE	GPU COMPUTE	HOST I/O	MASK AGREEMENT

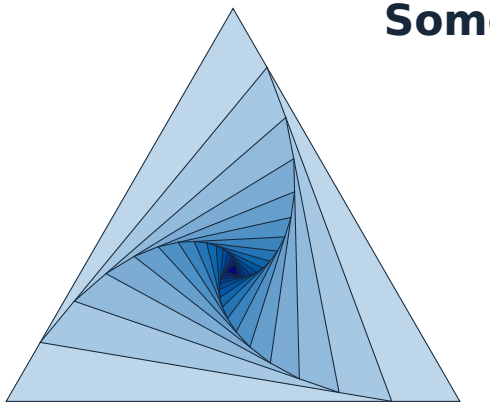

Some parametric custom figure

THE TAKEAWAY Use native LaTeX content within the benchmark layout.

Source: Theme demonstration · measured GPU benchmark results